

MCMC homework

Marko Medved

I. NUMERICAL INTEGRATION VS MONTE CARLO INTEGRATION

In this section we compare numerical integration to Monte Carlo integration on a simple domain where the standard numerical integration is possible. Concretely, we approximate the expected value of the (a, b, c)-PERT distribution:

$$p(x) = \frac{(x-a)^{\alpha-1} \cdot (c-x)^{\beta-1}}{B(\alpha, \beta) \cdot (c-a)^{\alpha+\beta-1}}$$

with $a = 0$, $b = 10$ and $c = 100$.

A. Trapezoidal rule

We begin by approximating the expected value using the Trapezoidal rule. To determine how many function evaluations are required, we first computed the exact expected value analytically: $\mathbb{E}[X] = 23.3$. We then applied the numerical method, incrementally increasing the number of function evaluations until the result was accurate to four decimal places (the difference between estimated and true value was less than 5×10^{-5}). This level of precision was achieved with 271 evaluations.

B. CLT estimation for Monte-Carlo

In this section, we estimate the expected value using Monte Carlo integration. Our objective is to determine how many random samples are needed to approximate the integral within two decimal places, with 95% confidence.

The Monte-Carlo approximation in this example is done by drawing independent samples $x_1, x_2, \dots, x_n \sim \mathcal{U}(a, c) = \mathcal{U}(0, 100)$, and estimating:

$$\mathbb{E}[X] = \int_a^c x \cdot p(x) dx \approx \frac{c-a}{n} \sum_{i=1}^n x_i p(x_i).$$

Let this estimator be denoted by $\hat{\mu}_n$. By the Central Limit Theorem, the standard error of $\hat{\mu}_n$ is approximately:

$$SE = \frac{\hat{\sigma}}{\sqrt{n}},$$

where $\hat{\sigma}^2$ is the sample variance of the integrand values $x_i p(x_i)$.

To estimate $\mathbb{E}[X]$ to 2 decimal places with 95% confidence, we require:

$$1.96 \cdot \frac{\hat{\sigma}}{\sqrt{n}} \leq 0.005.$$

So to estimate the number of points, we need to compute σ . Since this is the variance of our integrand, we decided to estimate it using a pilot run of 10^8 samples. Our estimation was $\hat{\sigma} = 19.08$ and then we calculated: $n \approx 55935632$.

C. Numerical samples verification

To verify our estimate of the required number of samples, we perform 200 independent Monte Carlo simulations. We then compute the difference between each estimate and the true expected value, and visualize the distribution of these differences as a density plot in Figure 1. We can see that almost all samples, as expected since we defined the 95% confidence, have an absolute difference to the true value less than 0.005. We can also see that the density is gathered around 0.

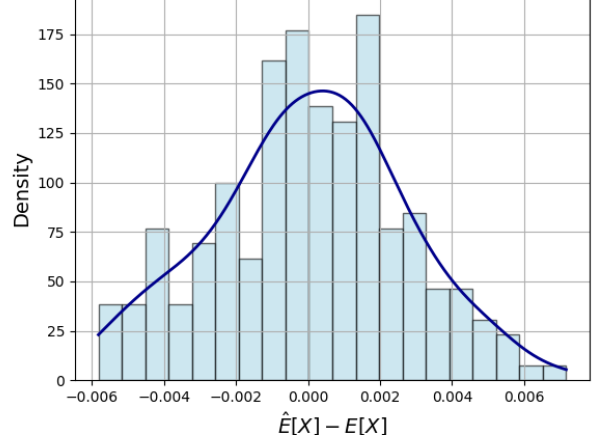


Figure 1. Density of differences of the estimated and true expected value

D. Comparison of results from both methods

As expected, the trapezoidal rule significantly outperformed the Monte Carlo method. It provided a deterministic result accurate to four decimal places using relatively few evaluations, while the Monte Carlo estimate required many more samples and achieved only two-decimal-place accuracy with 95% confidence. This highlights that for low-dimensional problems with simple domains, standard numerical integration methods are preferable to Monte Carlo methods.

II. IMPORTANCE SAMPLING

In this section we implement importance sampling for Monte Carlo integration to be able to compute the provided integral: $I = \int_0^1 x^{-3/4} \cdot e^{-x} dx$.

We plot the integrand ($f(x) = x^{-3/4} \cdot e^{-x}$) on Figure2. We can see that the bulk of the integral is gathered really close to 0.

A. Sampling from the Uniform distribution

Here we use the uniform distribution as the distribution we want to sample from. Since our uniform distribution is from 0 to 1, the estimation of the integral is simply:

$$I \approx \frac{1}{n} \sum_{i=1}^n f(x_i)$$

As we have seen on Figure 2, the function has a shape that is very concentrated at 0 so the uniform choice might not be the best.

B. Sampling from the provided distribution

Next we use the provided distribution ($q(x) = c \times x^{-3/4}$).

Firstly, we determine the c constant to be $1/4$ so that the distribution integrates to 1 over the domain $[0, 1]$.




Figure 2. Graph of the given integrand

Next, to implement inversion sampling, we need the CDF and its inverse. We calculated these functions to be:

$$F(x) = x^{1/4}$$

$$F^{-1}(x) = x^4$$

Then we simply sampled from the uniform distribution on the $[0, 1]$ interval, and transformed those samples with our inverse CDF to get the samples of $q(x)$.

Lastly, we implemented the importance sampling. Our integral can now be approximated as:

$$I \approx \frac{1}{n} \sum_{i=1}^n \frac{f(x_i)}{q(x_i)}$$

C. Comparison of results from both methods

Here we compare the results of two Monte Carlo integration methods. For each method, we drew 10 independent samples of size $n = 10^7$, computed the average estimate, and calculated the corresponding standard deviation. The results are summarized in Table I.

Sampling Distribution	Estimate	Standard Deviation
Uniform	3.3640	0.0890
$q(x)$	3.3795	0.0001

Table I
COMPARISON OF IMPORTANCE SAMPLING ESTIMATES

As shown in the table, using $q(x)$ as a proposal distribution results in a significantly lower standard deviation compared to uniform sampling. This improvement is expected, as $q(x)$ closely matches the shape of the integrand, leading to more samples being drawn from regions that contribute most to the integral. Consequently, the variance of the estimator is greatly reduced.

III. METROPOLIS-HASTINGS ALGORITHM

Here, we implement the Metropolis-Hastings (MH) algorithm to estimate the mean and variance of the parameters α and η of the Weibull distribution, which models our random variable.

Firstly, we compute the posterior from the given likelihood and prior:

$$p(\alpha, \eta | x) \propto \left(\prod_{i=1}^n \alpha \eta x_i^{\alpha-1} e^{-\eta x_i^\alpha} \right) \cdot e^{-\alpha-2\eta} \eta$$

Note that we use the data provided in the instructions. Note also that for all of the constructed chains we use the initial point of $[1, 1]$.

A. Multivariate normal proposal

To sample from the posterior distribution

$$p(\alpha, \eta | \vec{x}),$$

we firstly implemented the algorithm using a multivariate normal proposal distribution.

We tested several covariance matrices by running 5 independent MH chains for each one.

For each setting, we computed the sum of effective sample sizes (ESS) across all chains for each parameter (α and η).

The following covariance matrices were selected for testing:

- $\begin{bmatrix} 0.1 & 0.0 \\ 0.0 & 0.1 \end{bmatrix}$ Small step sizes in both directions
- $\begin{bmatrix} 1.0 & 0.0 \\ 0.0 & 1.0 \end{bmatrix}$ Large step sizes in both directions
- $\begin{bmatrix} 0.5 & 0.01 \\ 0.01 & 0.5 \end{bmatrix}$ Moderate scale with weak correlation
- $\begin{bmatrix} 0.5 & 0.05 \\ 0.05 & 0.5 \end{bmatrix}$ Moderate scale with stronger correlation
- $\begin{bmatrix} 1.0 & 0.0 \\ 0.0 & 0.1 \end{bmatrix}$ Larger steps for α , smaller steps for η
- $\begin{bmatrix} 0.1 & 0.0 \\ 0.0 & 1.0 \end{bmatrix}$ Smaller steps for α , larger steps for η

In Table II, we observe the effective sample sizes (ESS) obtained using different proposal covariance matrices.

Larger step sizes generally produced better results, as evidenced by the higher ESS values for both α and η . Interestingly, the highest ESS for each individual parameter was achieved using imbalanced covariance matrices (the last two entries). Specifically, the matrix with a large variance for α and small for η yielded the best ESS for α , while the reverse configuration performed best for η .

However, the most balanced performance across both parameters was achieved when using a covariance matrix with large variances in both directions:

$$\begin{bmatrix} 1.0 & 0.0 \\ 0.0 & 1.0 \end{bmatrix}$$

This choice provides a good trade-off, achieving high ESS for both parameters without favoring one over the other. That is why this covariance matrix will be used for further evaluation.

Covariance Matrix	ESS (α)	ESS (η)
$\begin{bmatrix} 0.1 & 0.0 \\ 0.0 & 0.1 \end{bmatrix}$	163.5	130.6
$\begin{bmatrix} 1.0 & 0.0 \\ 0.0 & 1.0 \end{bmatrix}$	622.9	549.5
$\begin{bmatrix} 0.5 & 0.01 \\ 0.01 & 0.5 \end{bmatrix}$	406.3	510.2
$\begin{bmatrix} 0.5 & 0.05 \\ 0.05 & 0.5 \end{bmatrix}$	418.4	378.0
$\begin{bmatrix} 1.0 & 0.0 \\ 0.0 & 0.1 \end{bmatrix}$	676.9	83.3
$\begin{bmatrix} 0.1 & 0.0 \\ 0.0 & 1.0 \end{bmatrix}$	194.7	594.9

Table II

EFFECTIVE SAMPLE SIZE (ESS) FOR α AND η ACROSS DIFFERENT PROPOSAL COVARIANCE MATRICES.

Then with the chosen matrix, we calculated the mean and the variance of the posteriors of α and η . The results are summarized as follows:

- Overall mean of α : 1.7010
- Variance of α : 0.3869
- Overall mean of η : 1.8327
- Variance of η : 0.6918

These statistics summarize the central tendency and dispersion of the parameter estimates obtained from the MCMC samples across all chains.

Next, we applied the other standard MCMC diagnostics. In Figure 3, we can see the traceplots for both parameters across all independent chains. The chains appear to mix well. There are no apparent trends, which indicates that stationarity has been reached.

Next, in Figure 4, we present the autocovariance plots for both parameters, using the first chain as a representative example. We observe that autocovariance remains relatively high for small lags, particularly for η , indicating a certain degree of dependence between consecutive samples. This behavior is consistent with the corresponding effective sample sizes reported in Table II, where lower ESS values reflect the reduced efficiency caused by this correlation.




Figure 3. Trace plot for the multivariate normal proposal




Figure 4. Autocovariance plot for the multivariate normal proposal

B. The provided proposal

Here we used:

$$q(\alpha', \eta' | \alpha, \eta) = \frac{1}{\alpha\eta} \exp\left(-\frac{\alpha'}{\alpha} - \frac{\eta'}{\eta}\right)$$

as our proposal. This is a product of two exponential distributions, so we needed to sample like this:

$$\alpha' \sim \text{Exp}(\lambda = 1/\alpha), \quad \eta' \sim \text{Exp}(\lambda = 1/\eta).$$

Then with the chosen proposal, we calculated the overall mean and variance of the parameters across all chains:

- α : mean = 1.6501, variance = 0.5179
- η : mean = 1.8212, variance = 0.8255

Next, we present the trace plot in Figure 5. It is evident that many proposals were rejected when using this exponential proposal distribution, as indicated by the flat segments in the traces.

In contrast, the autocovariance plot in Figure 6 shows that the autocovariance decreases slightly faster compared to the previous (normal) proposal distribution, suggesting improved mixing behavior.

Finally, we consider the Effective Sample Size (ESS) across all chains:

- Total ESS for α : 486.41
- Total ESS for η : 484.12

We can see that the ESS for both parameters is quite a lot higher in comparison to the multivariate normal distribution.

C. Comparison of both algorithms

Based on the results, we observe that the Multivariate Normal proposal outperformed the Exponential proposal in terms of effective sample size and overall sampling efficiency. The MVN proposal yielded higher ESS values for both α and η , indicating better mixing and more informative samples across chains. However, achieving this performance required careful tuning of the proposal covariance matrix, which adds complexity and additional effort to the implementation. In contrast, the Exponential proposal, while slightly less efficient, had the advantage of not requiring such tuning, making it simpler and more straightforward to apply in practice.

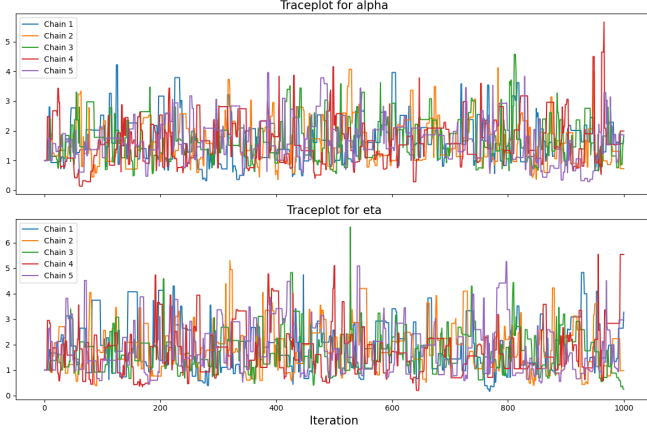


Figure 5. Trace plot for the exponential proposal

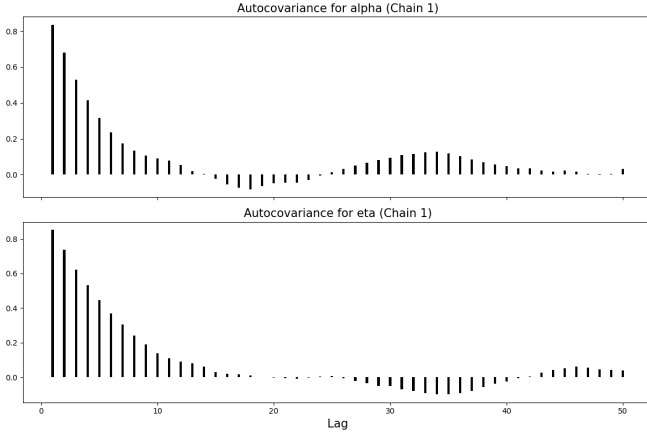


Figure 6. Autocovariance plot for the exponential proposal

D. Probability estimation

Lastly, we computed the probability $P((\alpha, \eta) \in [1.3, \infty) \times [1.3, \infty))$ using samples from both proposals. The estimated probabilities were:

- Multivariate Normal proposal: $\hat{P} = 0.5562$
- Exponential proposal: $\hat{P} = 0.4635$

These results indicate that the multivariate normal proposal yields a higher estimated probability in the specified region compared to the exponential proposal.