

MCMC homework

Marko Medved

I. NUMERICAL INTEGRATION VS MONTE CARLO INTEGRATION

In this section we compare numerical integration to Monte Carlo integration on a simple domain where the standard numerical integration is possible. Concretely, we approximate the expected value of the (a, b, c)-PERT distribution:

$$p(x) = \frac{(x-a)^{\alpha-1} \cdot (c-x)^{\beta-1}}{B(\alpha, \beta) \cdot (c-a)^{\alpha+\beta-1}}$$

with $a = 0$, $b = 10$ and $c = 100$.

A. Trapezoidal rule

We begin by approximating the expected value using the Trapezoidal rule. To determine how many function evaluations are required, we first computed the exact expected value analytically: $\mathbb{E}[X] = 23.3$. We then applied the numerical method, incrementally increasing the number of function evaluations until the result was accurate to four decimal places (the difference between estimated and true value was less than 5×10^{-5}). This level of precision was achieved with 271 evaluations.

B. CLT estimation for Monte-Carlo

In this section, we estimate the expected value using Monte Carlo integration. Our objective is to determine how many random samples are needed to approximate the integral within two decimal places, with 95% confidence.

The Monte-Carlo approximation in this example is done by drawing independent samples $x_1, x_2, \dots, x_n \sim \mathcal{U}(a, c) = \mathcal{U}(0, 100)$, and estimating:

$$\mathbb{E}[X] = \int_a^c x \cdot p(x) dx \approx \frac{c-a}{n} \sum_{i=1}^n x_i p(x_i).$$

Let this estimator be denoted by $\hat{\mu}_n$. By the Central Limit Theorem, the standard error of $\hat{\mu}_n$ is approximately:

$$SE = \frac{\hat{\sigma}}{\sqrt{n}},$$

where $\hat{\sigma}^2$ is the sample variance of the integrand values $x_i p(x_i)$.

To estimate $\mathbb{E}[X]$ to 2 decimal places with 95% confidence, we require:

$$1.96 \cdot \frac{\hat{\sigma}}{\sqrt{n}} \leq 0.005.$$

So to estimate the number of points, we need to compute σ . Since this is the variance of our integrand, we decided to estimate it using a pilot run of 10^8 samples. Our estimation was $\hat{\sigma} = 19.08$ and then we calculated: $n \approx 55935632$.

C. Numerical samples verification

To verify our estimate of the required number of samples, we perform 200 independent Monte Carlo simulations. We then compute the difference between each estimate and the true expected value, and visualize the distribution of these differences as a density plot in Figure 1. We can see that almost all samples, as expected since we defined the 95% confidence, have an absolute difference to the true value less than 0.005. We can also see that the density is gathered around 0.

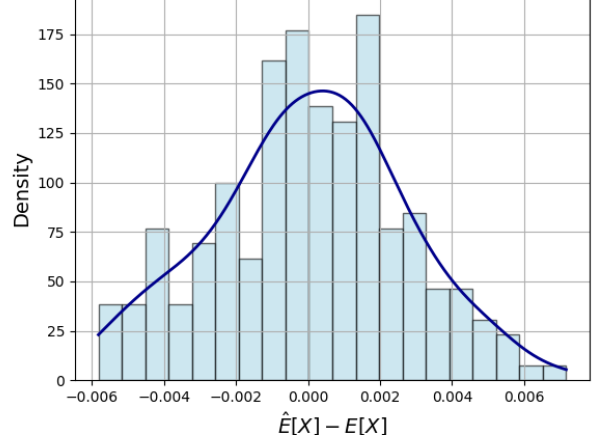


Figure 1. Density of differences of the estimated and true expected value

D. Comparison of results from both methods

As expected, the trapezoidal rule significantly outperformed the Monte Carlo method. It provided a deterministic result accurate to four decimal places using relatively few evaluations, while the Monte Carlo estimate required many more samples and achieved only two-decimal-place accuracy with 95% confidence. This highlights that for low-dimensional problems with simple domains, standard numerical integration methods are preferable to Monte Carlo methods.

II. IMPORTANCE SAMPLING

In this section we implement importance sampling for Monte Carlo integration on the provided integral: $I = \int_0^1 x^{-3/4} \cdot e^{-x} dx$

We plot the integrand ($f(x) = x^{-3/4} \cdot e^{-x}$) on Figure 2. We can see that the bulk of the integral is gathered really close to 0.

A. Uniform distribution

Here we use the uniform distribution as the distribution we want to sample from. Since our uniform distribution is from 0 to 1, the estimation of the integral is simply:

$$I \approx \frac{1}{n} \sum_{i=1}^n f(x_i)$$

As we have seen on Figure 2, the function has a shape that is very concentrated at 0 so the uniform choice might not be the best.

B. The provided distribution

Next we use the provided distribution ($q(x) = c \times x^{-3/4}$).

Firstly, we determine the c constant to be $1/4$ so that the distribution integrates to 1 over the domain $[0, 1]$.




Figure 2. Graph of the given function

Next, to implement inversion sampling, we need the CDF and it's inverse. We calculated these functions to be:

$$F(x) = x^{1/4}$$

$$F^{-1}(x) = x^4$$

Then we simply sampled from the uniform distribution on the $[0, 1]$ interval, and transformed those samples with our inverse CDF to get the samples of $q(x)$.

Lastly, we implemented the importance sampling. Our integral can now be approximated as:

$$I \approx \frac{1}{n} \sum_{i=1}^n \frac{f(x_i)}{q(x_i)}$$

C. Comparison of results from both methods

Here we compare the results of two Monte Carlo integration methods. For each method, we drew 10 independent samples of $size = 10^7$, computed the average estimate, and calculated the corresponding standard deviation. The results are summarized in Table I.

Sampling Distribution	Estimate	Standard Deviation
Uniform	3.3640	0.0890
$q(x)$	3.3795	0.0001

Table I

COMPARISON OF IMPORTANCE SAMPLING ESTIMATES

As shown in the table, using $q(x)$ as a proposal distribution results in a significantly lower standard deviation compared to uniform sampling. This improvement is expected, as $q(x)$ closely matches the shape of the integrand, leading to more samples being drawn from regions that contribute most to the integral. Consequently, the variance of the estimator is greatly reduced.

III. METROPOLIS-HASTINGS ALGORITHM

Here, we implement the Metropolis-Hastings algorithm to compute the mean and the variance of the η and α , which are the parameters of the Weibull distribution which is the distribution of our random variable.

Firstly, we compute the posterior from the given likelihood and prior:

$$p(\alpha, \eta | x) \propto \left(\prod_{i=1}^n \alpha \eta x_i^{\alpha-1} e^{-\eta x_i^\alpha} \right) \cdot e^{-\alpha-2\eta} \eta$$